

About the job

Your new company

A scaling Technology Group, based in the heart of Brisbane CBD is seeking to hire a top-tier AI Engineer, to be embedded within their talent team. Providing 'tech for good' technology, this business has been built on cutting-edge technology and continues to push the boundaries of Gen AI, as they grow their business, and delight their customers. This is a truly unique opportunity to be part of something special, and join a growing technology business.

Your new role

Our customer is seeking a full-time Generative AI Engineer to join their Brisbane team. In this role, you will be responsible for developing and implementing AI solutions that enhance their platform. The key responsibilities include conducting research and development of AI/ML solutions to enhance platform capabilities, fine-tuning and optimising Large Language Models (LLMs) for domain-specific applications, and developing LLM-based applications to automate and streamline complex workflows. You will also perform prompt engineering and optimisation to ensure high-quality AI outputs, integrate AI solutions into production environments with scalability and reliability in mind, and design, develop, and maintain APIs for seamless AI integration. This is a role that will require plenty of collaboration with a talented peer group, and the chance to bring your own ideas to the emerging space.

What you'll need to succeed

You will need over 2 years of experience in developing and deploying AI/ML solutions or similar technologies, and at least 1 year of experience in deploying AI applications to production environments. The ideal candidate should have expertise in building applications using Large Language Models (LLMs) and experience with LLM frameworks such as LangChain and LlamaIndex. Proficiency with vector databases and LLM Ops for model management is essential, along with strong experience in API integration and design. Familiarity with the Azure technology stack, including Azure Machine Learning and Azure Cognitive Services, is also required. Additionally, a solid background in Natural Language Processing (NLP), deep learning, and classical machine learning is necessary.

What you'll get in return

You will get the chance to work on cutting edge technology, as part of a talented peer group, and be exposed to future Gen AI trends and technology. Hybrid environment, with a couple of days in the office and the rest at home. Salary is highly dependent on experience.

What you need to do now

If you're interested in this role, click 'apply now' to forward an up-to-date copy of your CV, or call us now.

If this job isn't quite right for you, but you are looking for a new position, please contact us for a confidential discussion about your career.